

Operating Systems

Recitation 5

Plan

- Modes of operations
- What causes them
- Interrupts & Exceptions

Need protection

- Kernel privileged – cannot trust user processes
 - User processes may be malicious or buggy
- Must protect
 - User processes from one another
 - Kernel from user processes




Hardware assisted memory protection

- **Paging** – we discussed virtual memory...
 - Every process has its own virtual memory address space
 - OS uses tables for translation, etc.
 - Mapped to physical memory
 - We occasionally swap to disk
 - ***P1*** can't access memory of ***P2***


Segmentation

- Briefly mentioned too
 - Kernel code and data mapped in every process to same segments
 - User code etc. in different segments



Dual mode of operation

- **Hardware** (again) *to the rescue!*
- Privileged (+ non-privileged) operations in kernel mode
- Only non-privileged operations in user mode



Segmentation + modes operation = ?

- Don't want meddling user processes to mess with kernel stuff
- Each segment has **Descriptor Privilege Level (DPL)**
 - kernel segments: 0
 - user segments: 3
- Each thread has **Current Privilege Level (CPL)**
- Can only access segments when **CPL ≤ DPL**
- CPL controlled by OS (when switching modes) and checked against each segment's DPL

Scheduling

- Timer interrupt (scheduler)
- Kernel periodically gets control back



Event driven mode-switch

- Events cause mode-switch
 - Interrupts: raised by devices to get OS attention
 - System calls: issued by user processes to request system services
 - Exceptions: illegal instructions (e.g., division by 0)

Why do we need interrupts?

- People like connecting devices
 - A computer is much more than the CPU
 - Keyboard, mouse, screen, disk drives...
 - These devices occasionally need CPU service
 - But we can't predict when...
 - External events typically occur on a macroscopic timescale
 - We want to keep the CPU busy between events
- ⇒ **Need a way for CPU to find out devices need attention**

Interrupts

Electric signal to processor, signifies event that requires immediate attention

- Causes kernel to stop performing current code and use “special code” to handle interrupt
- Mostly asynchronous and raised by devices to get OS attention
 - Keyboard keystroke
- **/proc/interrupts** - lists all hardware devices and how many interrupts received by each on each CPU (more in /proc/pci)
- In x86 (32bit) 256 types of interrupts

Interrupt Handler

- Often does very little at the expense of current running process
- Defers whatever it can to later execution by some kernel thread
- Example: Networking
 - time-critical work - copy network data packet off network card, respond to card and continue
 - deferred work - process data packet, pass to correct application (browser,...)

Interrupt view of CPU

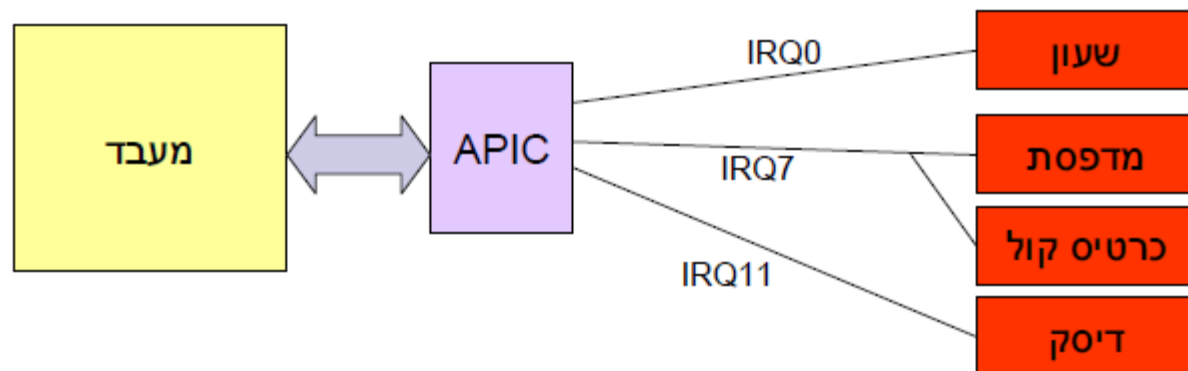
```
while (fetch next instruction) {  
    run instruction;  
  
    if (there is an interrupt) {  
        process interrupt  
    }  
}
```

A deeper look

```
while (fetch next instruction) {  
    run instruction;  
    if (there is an interrupt) {  
        switch to process kernel stack if necessary  
        save CPU context and error code if any  
        find OS-provided interrupt handler  
        jump to handler  
        restore CPU context when handler returns  
    }  
}
```


Finding the interrupt handler

- Processor has **[A]PIC**
[Advanced] Programmable Interrupt Controller
 - In multicore systems - sophisticated OS mechanisms, frontend APIC routes to core-specific local APIC. Routing done with masking...
 - `/proc/irq/XXX/smp_affinity`



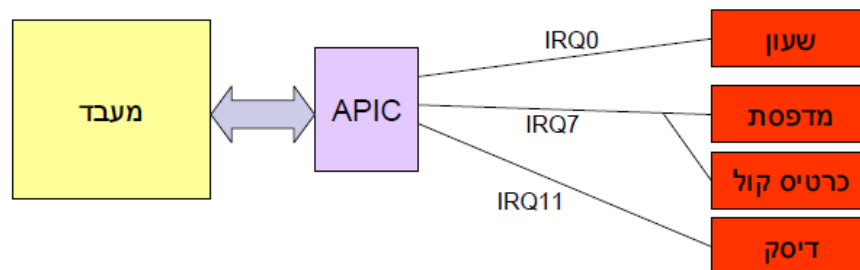
APIC

- On interrupt device sends signal to IRQ (Interrupt ReQuest) line
- IRQ sharing – need to check all devices that might have caused it
- Exceptions (sync. interrupts) do not use APIC




APIC

- APIC notifies CPU about specific interrupt and holds
 - Doesn't serve interrupts meanwhile
- Linux ACKs specific interrupt
 - Allows APIC to receive more interrupts from other types
- Later when done, OS informs APIC it is ready to handle interrupts from specific type too

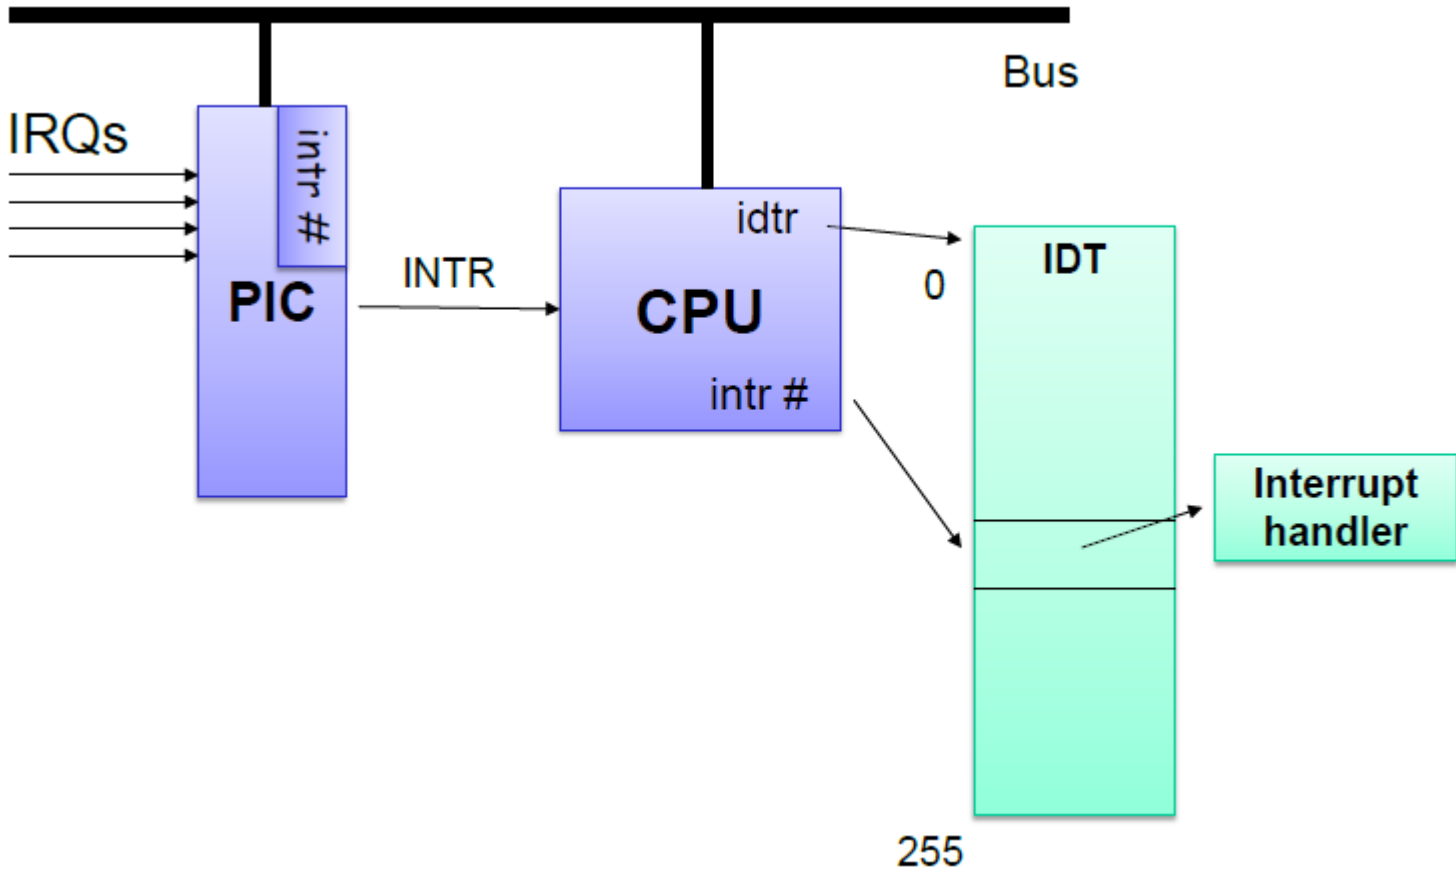


Interrupt Descriptor Table

- Hardware maps interrupt type to interrupt number
- On boot OS sets up “interrupt vector” or Interrupt Descriptor Table (IDT) in memory
 - Each entry is an interrupt handler
- OS lets hardware know IDT base in *idtr* register
- Hardware finds handler using interrupt num: `IDT[intr_number]`



All together now



IDT entries

- 0: divide by 0
- 1: debug (for single stepping)
- 3: breakpoint
- 14: page fault
- [128 used for syscall, soon more details...]
- 251-255 inter-processor interrupts

Interrupt Request

`unsigned int do_IRQ(struct pt_regs regs);`

- C function
- In kernel 2.6 `arch/i386/kernel/irq.c`
- Find relevant interrupt (from APIC registers)
- Invoke relevant handler called **Interrupt Service Routine (ISR)**
- ACK to APIC on relevant IRQ
 - Can handle more such interrupts

Interrupt Service Routine (ISR)

- Handles interrupts for specific device
 - Read keystroke code
 - Read data from network card
 - Perform I/O from disk
- Implemented by device drivers

Why not polling?

- CPU *periodically checks each device* to see if it needs service

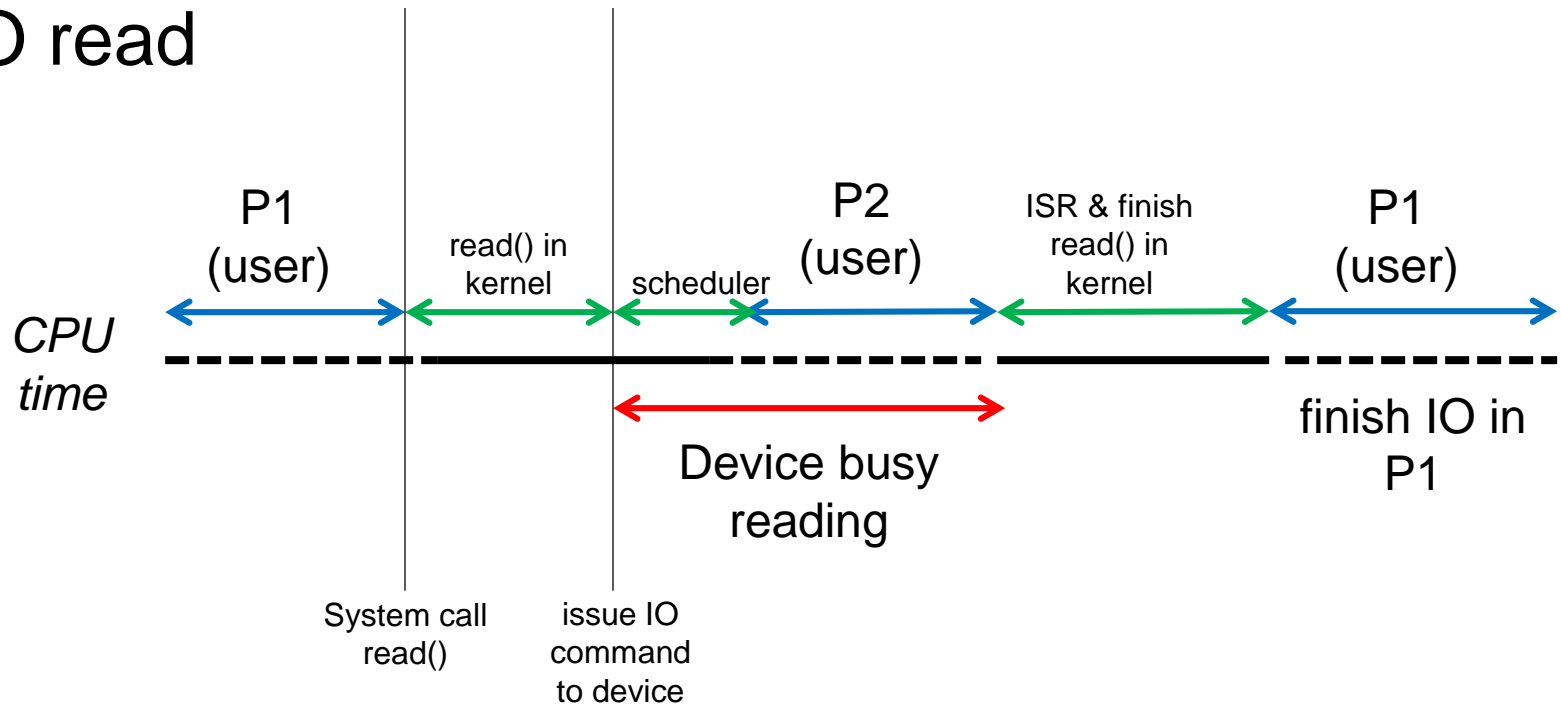
“Polling is like picking up your phone every few seconds to see if you have a call. ...”

- Takes CPU time when no requests are pending



Why not polling?

- IO read

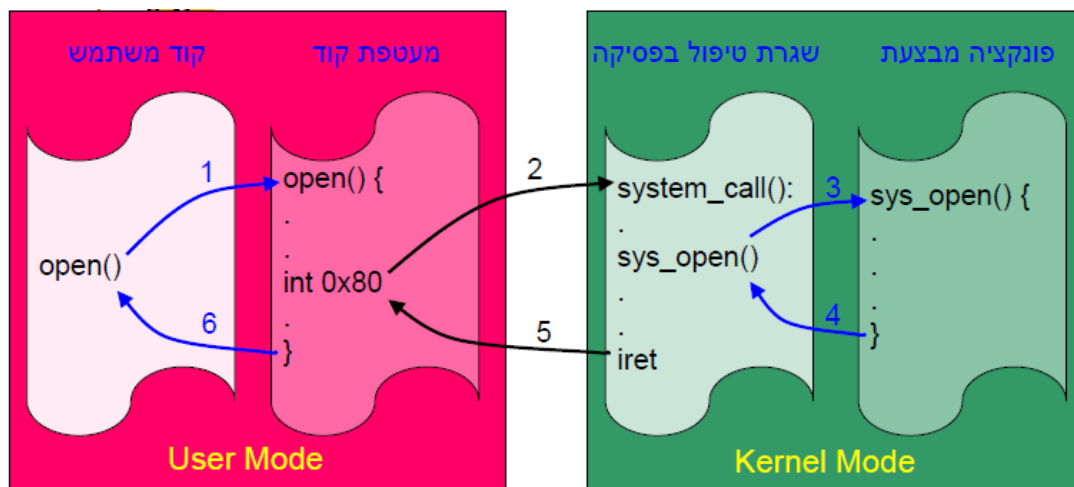


Exceptions == sync. interrupts

- Issued by processor
- *Exception in program execution*
 - Faults; offending instruction is retried
 - Division by 0
 - Traps; deliberate in code. instruction is not retried
 - Usually for debugging
 - Aborts; major error
 - hardware failure
- In multicore, cores often communicate via interrupts
 - And each core can handle different interrupts...

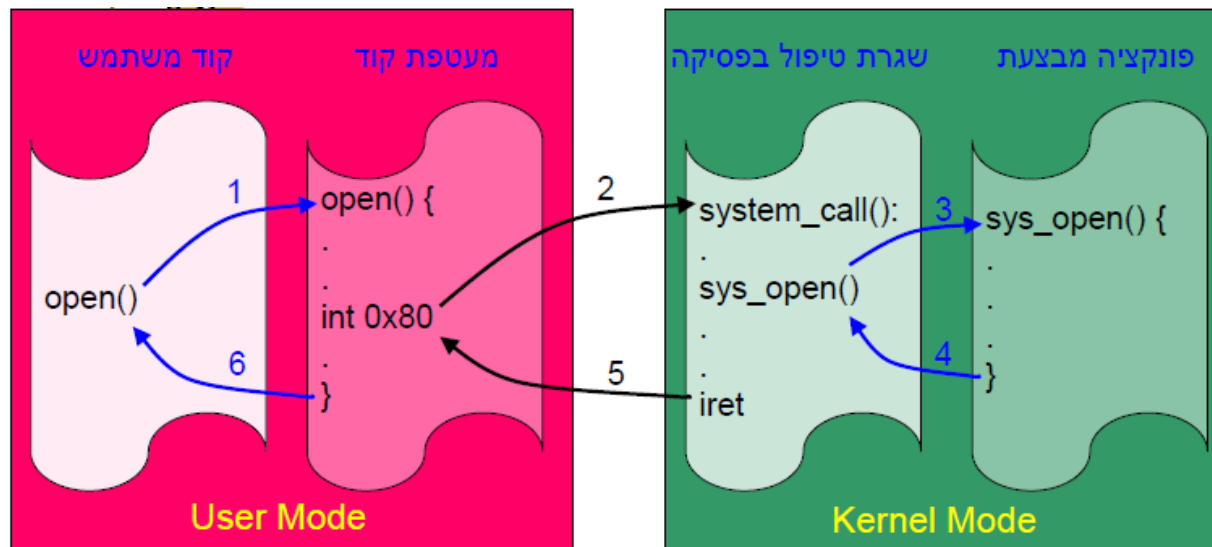
Programmed exceptions == software interrupts

- Mechanisms for special exceptions likely to occur
- Used for system calls! (and debugging)
 - Interrupt 128 (int 0x80) which invokes routine `system_call()`
 - Windows: int 0x2e



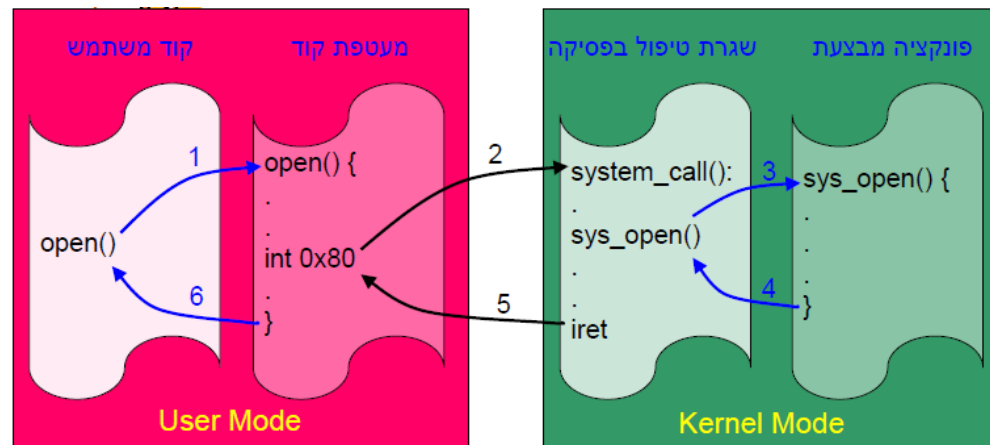
Programmed Exceptions

- On boot system call table (vector) initialized in kernel
 - Maps system call number to relevant implementation
- On demand user process sets up system call number & arguments in designated registers
- User process causes interrupt 128



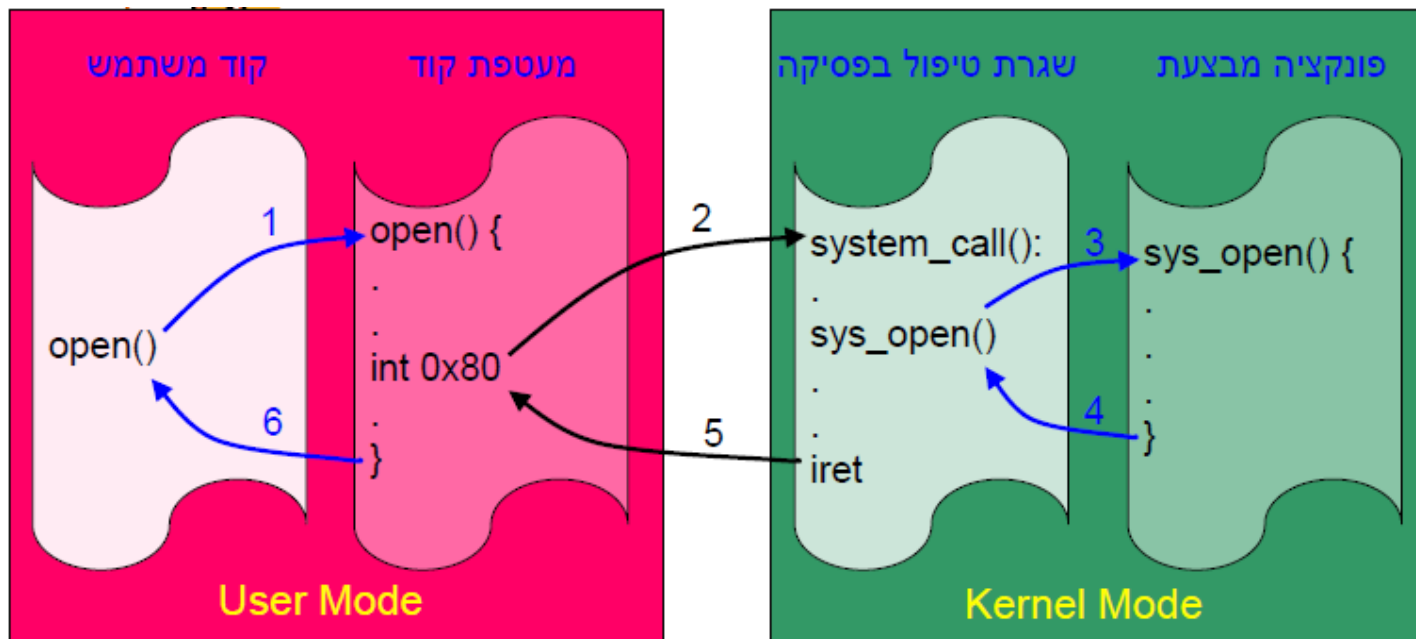
Programmed Exceptions

- Hardware switches to kernel mode and invokes relevant interrupt handler `system_call()`
- Kernel looks up syscall table using system call number from relevant register
- Kernel invokes the corresponding function
- Kernel returns to original execution point



Programmed Exceptions

- How is it used for debugging?



Linux Signals

Inform process async. on some event

- Form of IPC for messages from OS about events
 - 32 types in UNIX
- Some signals are handled automatically
 - KILL, STOP
- Others handled by signal handler
 - OS queues signals (only 1 for every type!)
 - Upon returning to user-mode process handles them using handler

Exception handling in Linux

- Reminder - exceptions are sync. Interrupts (faults, traps,...)
- Exception handler - saves relevant data on stack
- Then sends signal (async event...) to relevant process
- Signal type specific to exception
 - Divide by zero will send **SIGFPE**
 - Ctrl + C sends **SIGINT**
- Before returning to user mode, process checks signal handler

Linux Signals

Some signals are process-wide (KILL...)

- Threads can handle signals separately (using signal masks, somewhat complex)
- Want to handle some signals on your own? Implement your own signal handler...

Signals vs. interrupts

- Who's talking
 - Interrupts → communication between CPU and OS kernel
 - Signals → communication between the OS kernel and OS processes (and themselves)
- Assumption
 - Interrupts assume OS is frozen and waiting for them to finish
 - Signals obviously do not